

Motif finding with Gibbs sampling

CS 466

Saurabh Sinha

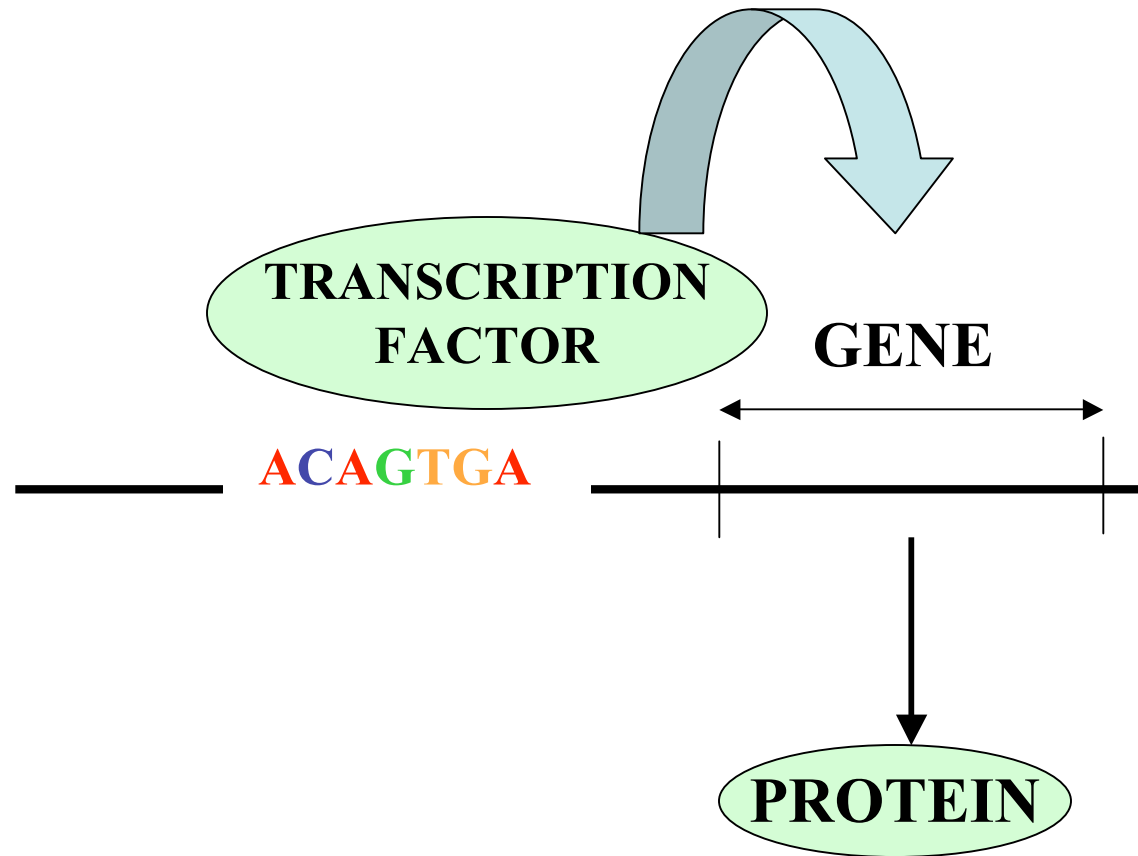
Regulatory networks

- Genes are **switches**, transcription factors are (one type of) input **signals**, proteins are **outputs**
- Proteins (**outputs**) may be transcription factors and hence become **signals** for other genes (**switches**)
- This may be the reason why humans have so few genes (the circuit, not the number of switches, carries the complexity)


Decoding the regulatory network

- Find patterns (“motifs”) in DNA sequence that occur more often than expected by chance
 - These are likely to be binding sites for transcription factors
 - Knowing these can tell us if a gene is regulated by a transcription factor (i.e., the “switch”)

Transcriptional regulation




Transcriptional regulation



A motif model

- To define a motif, let's say we know where the motif starts in the sequence
- The motif start positions in their sequences can be represented as $\mathbf{s} = (s_1, s_2, s_3, \dots, s_t)$



Motifs: Matrices and Consensus

Alignment

a	G	g	t	a	c	T	t
C	c	A	t	a	c	g	t
a	c	g	t	T	A	g	t
a	c	g	t	C	c	A	t
C	c	g	t	a	c	g	G

Matrix

A	3	0	1	0	3	1	1	0
C	2	4	0	0	1	4	0	0
G	0	1	4	0	0	0	3	1
T	0	0	0	5	1	0	1	4

Consensus **A C G T A C G T**

- Line up the patterns by their start indexes

$$\mathbf{s} = (s_1, s_2, \dots, s_t)$$

- Construct “position weight matrix” with frequencies of each nucleotide in columns
- Consensus nucleotide in each position has the highest frequency in column

Position weight matrices

- Suppose there were t sequences to begin with
- Consider a column of a position weight matrix
- The column may be $(t, 0, 0, 0)$
 - A perfectly conserved column
- The column may be $(t/4, t/4, t/4, t/4)$
 - A completely uniform column
- “Good” profile matrices should have more conserved columns

Information Content

- In a PWM, convert frequencies to probabilities
- PWM W : $W_{\beta k}$ = frequency of base β at position k
- q_{β} = frequency of base β by chance
- Information content of W :

$$\sum_k \sum_{\beta \in \{A, C, G, T\}} W_{\beta k} \log \frac{W_{\beta k}}{q_{\beta}}$$

Information Content

- If $W_{\beta k}$ is always equal to q_{β} , i.e., if W is similar to random sequence, information content of W is 0.
- If W is different from q , information content is high.

Detecting Subtle Sequence Signals: a Gibbs Sampling Strategy for Multiple Alignment

Lawrence et al. 1993

Motif Finding Problem


- Given a set of sequences, find the motif shared by all or most sequences, while its starting position in each sequence is unknown
- Assumption:
 - Each motif appears exactly once in one sequence
 - The motif has fixed length

Generative Model

- Suppose the sequences are aligned, the aligned regions are generated from a motif model
- Motif model is a PWM. A PWM is a position-specific multinomial distribution.
 - For each position i , a multinomial distribution on (A,C,G,T):
 $q_{iA}, q_{iC}, q_{iG}, q_{iT}$
- The unaligned regions are generated from a background model: p_A, p_C, p_G, p_T

Notations

- Set of symbols: Σ
- Sequences: $S = \{S_1, S_2, \dots, S_N\}$
- Starting positions of motifs: $A = \{a_1, a_2, \dots, a_N\}$
- Motif model (θ) : $q_{ij} = P(\text{symbol at the } i\text{-th position} = j)$
- Background model: $p_j = P(\text{symbol} = j)$
- Count of symbols in each column: c_{ij} = count of symbol, j , in the i -th column in the aligned region



Probability of data given model

$$P(S \mid A, \theta) = \prod_{i=1}^W \prod_{j=1}^{|\Sigma|} q_{ij}^{c_{ij}} \quad P(S \mid A, \theta_0) = \prod_{i=1}^W \prod_{j=1}^{|\Sigma|} p_j^{c_{ij}}$$

Scoring Function

- Maximize the log-odds ratio:

$$P(S \mid A, \theta) = \prod_{i=1}^W \prod_{j=1}^{|\Sigma|} q_{ij}^{c_{ij}} \quad P(S \mid A, \theta_0) = \prod_{i=1}^W \prod_{j=1}^{|\Sigma|} p_j^{c_{ij}}$$

$$F = \log \frac{P(S \mid A, \theta)}{P(S \mid A, \theta_0)} = \sum_{i=1}^W \sum_{j=1}^{|\Sigma|} c_{ij} \log \frac{q_{ij}}{p_j}$$

- Is greater than zero if the data is a better match to the motif model than to the background model

Optimization and Sampling

- To maximize a function, $f(x)$:
 - Brute force method: try all possible x
 - Sample method: sample x from probability distribution: $p(x) \sim f(x)$
 - Idea: suppose x_{\max} is argmax of $f(x)$, then it is also argmax of $p(x)$, thus we have a high probability of selecting x_{\max}

Markov Chain sampling

- To sample from a probability distribution $p(x)$, we set up a Markov chain s.t. each state represents a value of x and for any two states, x and y , the transitional probabilities satisfy:

$$p(x)\Pr(x \rightarrow y) = p(y)\Pr(y \rightarrow x)$$

- This would then imply that if the Markov chain is “run” for “long enough”, the probability thereafter of being in state x will be $p(x)$

$$\lim_{N \rightarrow \infty} \frac{1}{N} C(x) = p(x)$$

Gibbs sampling to maximize F

- Gibbs sampling is a special type of Markov chain sampling algorithm
- Our goal is to find the optimal $A = (a_1, \dots, a_N)$
- The Markov chain we construct will only have transitions from A to alignments A' that differ from A in only one of the a_i
- In round-robin order, pick one of the a_i to replace
- Consider all A' formed by replacing a_i with some other starting position a_i' in sequence S_i
- Move to one of these A' probabilistically
- Iterate the last three steps

Algorithm

Randomly initialize A^0 ;

Repeat:


- (1) randomly choose a sequence z from S ;
 $A^* = A^t \setminus a_z$; compute θ^t from A^* ; $q_{ij} = \frac{c_{ij}}{\sum_k c_{ik}}$
- (2) sample a_z according to $P(a_z = x)$, which is proportional to Q_x/P_x ; update $A^{t+1} = A^* \cup x$;

Select A^t that maximizes F ;

Q_x : the probability of generating x according to θ^t ;


P_x : the probability of generating x according to the background model

Algorithm




Current solution A^t

Algorithm




Choose one a_z to replace

Algorithm




For each candidate site x in sequence z ,
calculate Q_x and P_x :
Probabilities of sampling x from motif model and background model resp.

Algorithm




Among all possible candidates, choose one (say x) with probability proportional to Q_x/P_x

Algorithm



Set $A^{t+1} = A^* \cup x$

Algorithm



Repeat

Local optima

- The algorithm may not find the “global” or true maximum of the scoring function
- Once “ A^t ” contains many similar substrings, others matching these will be chosen with higher probability
- Algorithm will “get locked” into a “local optimum”
 - all neighbors have poorer scores, hence low chance of moving out of this solution